

# **Supercomputing Spin Up Part 2 - Job Submission**

September 5, 2024

Brandon Reyes



# Supercomputing Spin Up - Job Submission

## Instructor: Brandon Reyes

- Research Computing
- Website: [www.rc.colorado.edu](http://www.rc.colorado.edu)
- Helpdesk: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Slides:  
[https://github.com/ResearchComputing/supercomputing\\_spinup\\_job\\_submission\\_short\\_course](https://github.com/ResearchComputing/supercomputing_spinup_job_submission_short_course)
- Survey: <http://tinyurl.com/curc-survey18>



GitHub link for  
presentation



Link to survey



# Learning Objectives

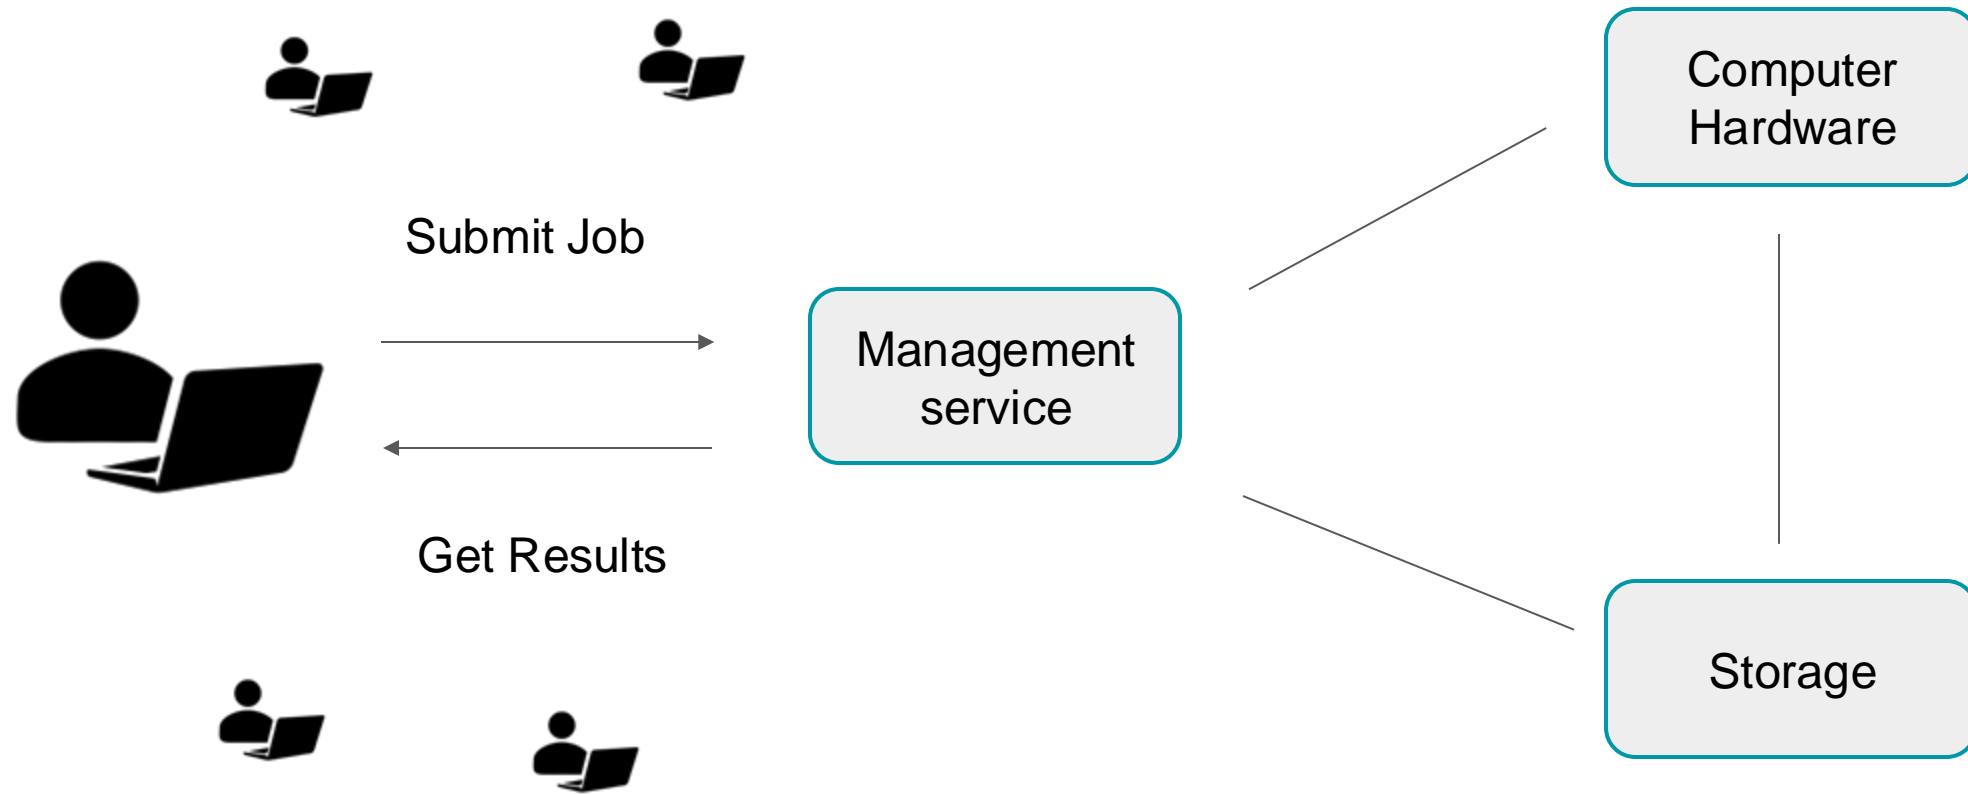
- Obtain a clear overview of job submission on an HPC system
- Learn about submitting both a batch and an interactive job to Alpine using the terminal

# Session Overview

- General overview of Job submission
- Introduction to batch jobs
  - Submission of a batch job
- Checking/monitoring jobs
- Utilizing software in a job
- Introduction to interactive jobs

Remember, when you login to the HPC system you are put on a login node. You need to then gain access to a compute node to run software.

# General Overview of Job Submission




# Introduction to Jobs

- Because our clusters are shared resources with many users trying to utilize available compute with their applications, we need a system to divide compute in a simple and fair system
- SLURM
  - **S**imple **L**inux **U**tility for **R**esource **M**anagement
- Through SLURM, users can grab allotments of compute resources called Jobs
- 2 Types of Jobs
  - **Batch Jobs**
  - **Interactive Jobs**



# General Overview of Job Submission



# Job directives

- Directives are special flags that specify what type of HPC resources you would like to use
  - “I want to run on 1 CPU for 1 hour”
- Used for both interactive and batch jobs
- Can be specified via the command line or in a job script
- All directives shown are specific to SLURM

# Common directives

- **Partition – A collection of compute nodes**
  - `--partition=<partition_name>`
- **Quality of service (QoS) – System defined constraints for a job (more on this later!)**
  - `--qos=<qos>`
- **Allocation – Account to “charge to”**
  - `--account=<account_name>`
- **Number of nodes to run on**
  - `--nodes=<nodes>`
- **Number of cores to run on**
  - `--ntasks=<number-of-tasks>`

# Common directives

- **Wall time** – How long you want to run on these resources
  - `--time=<wall time>`
- **Job name**
  - `--job-name=<jobname>`
- **Output** – Where all output that would be written to the terminal should go
  - `--output=<name>`
- **Send an email when events happen in the job**
  - `--mail-type=<type>`
- **Email address to send updates to**
  - `--mail-user=<user>`

# Alpine Partitions

- Partitions are a collection of compute nodes e.g. computers with common characteristics

Partition	Description	# of nodes	RAM/core (GB)	cores/node	GPUs/node
amilan	AMD Milan CPU node	347	3.8	64	0
ami100	AMD MI100 GPU node	8	3.8	64	3
aa100	Nvidia A100 GPU node	12	3.8	64	3
amem	High-memory node	22	21.5	48	0
atesting	Multi-node testing nodes	2	3.8	64	0
atesting_a100	Nvidia A100 testing node	1	3.8	64	3
atesting_mi100	AMD MI100 testing node	1	3.8	64	3

# Quality of Service (QoS)

- Quality of Service specifies additional constraints for a job
  - On Alpine, QoS can be used to run long jobs, specify testing partitions, and select high-memory nodes

QoS	Description	Max wall time	Max jobs/user	Max nodes/user
normal	Default QoS	24 H	1000	128
long	For jobs needing longer wall times	7 D	200	20
mem	High-memory jobs	7 D	n/a	12
testing	For jobs submitted to testing partitions	1 H	1	n/a

# Batch Jobs

- **Batch Jobs** are jobs you submit to the scheduler that are run later without supervision
  - By far the most common job on Alpine
  - Requires a job script
- A job script is simply a script that includes **SLURM directives** (resource specifics) ahead of any commands.

# Anatomy of a job script

- It's just a bash script with SLURM specific directives!

```
#!/bin/bash
```

```
## Directives
```

```
#SBATCH --<option>=<value>
```

```
## Software
```

```
module load <software>
```

```
## User scripting
```

```
<command>
```





# Directives in a job script

```
#SBATCH --<option>=<value>
```

# Example job script

```
#!/bin/bash

## Directives
#SBATCH --ntasks=1           # Number of requested tasks/cores
#SBATCH --time=0:01:00       # Max run time
#SBATCH --partition=amilan    # Specify Alpine CPU node
#SBATCH --output=test_%j.out  # Rename standard output file

## Software
module purge                  # Purge all existing modules

## User commands
echo "This is a test of user $USER"
```

# Submitting a Job script

- Once a job script has been constructed you must submit it to the HPC system using SLURM
  - Done using `sbatch`
- If we created the job script “my\_first\_job.sh” then we would submit it as follows:

```
sbatch /path/to/my_first_job.sh
```

# Job output

- Once a job completes its execution, the standard output of the script will be redirected to an output file.
  - Great for debugging!
  - Could be different from output generated by your application
  - File is created in directory job was run unless specified in your `--output` directive.
  - If the directive `--output` is not provided, then a generic file name will be used (slurm\_XXXXXX.out).

# Checking your jobs

- **squeue**: Monitor your jobs status **in queue and while running**:
  - By default, shows all jobs in queue can specify using:

```
$ squeue -u <username>  
$ squeue -p <partition>
```

- **sacct**: Check back on usage statistics of **previous Jobs**
  - By default, only checks all jobs from the start of the current day can specify using:

```
$ sacct -u <username>  
$ sacct --start=MM/DD/YY -u <username>  
$ sacct -j <job-id>
```

# Checking your jobs

- Another method of checking details of your job while running is with `scontrol`
  - Advanced command usually used by system administrators, but you can use it too!

```
$ scontrol show job <job number>
```

- To check the percentage of CPU and memory usage of a job **after it completes**, use `seff`

```
$ module load slurmttools  
$ seff <job number>
```

# Software and Jobs

- Okay so running a job is easy, but how do I run a job with my software?
- We can utilize all the software we discussed in the previous talk in the job script!
- Any non-GUI related commands you would run from the command line can be put in your job script!
  - If you would like to use GUI applications, you will need X11 forwarding and an interactive job

# Software job script example

```
#!/bin/bash

## Directives
#SBATCH --ntasks=1                # Number of requested tasks/cores
#SBATCH --time=0:01:00            # Max run time
#SBATCH --partition=amilan        # Specify Alpine CPU node
#SBATCH --output=test_%j.out      # Rename standard output file

## Software
module purge                      # Purge all existing modules
module load anaconda              # Load Anaconda
conda activate <my-conda-environment> # Activate CONDA environment

## Run Python script
python my_cool_script.py
```





# Interactive jobs

- Interactive jobs are used to gain access to a compute node in real time
  - Great for testing and debugging!
- Interactive jobs can be subject to normal wait times!
- We can get access to a compute node interactively with `sinteractive`
- We also have a specialized command called `acompile` that provides you with an interactive session
  - Access to quick resources
  - Limit of 4 CPUs for up to 12 hours

# Running an interactive job

- Here we will run a Python script using an interactive job
- First request resources:

```
$ sinteractive --partition=atesting --qos=testing --time=00:10:00 --ntasks=1
```

- When the job starts you will be put on a compute node and you can execute your commands:

```
$ module load anaconda  
$ python my_very_cool_script.py
```

- To quit:

```
$ exit
```



GitHub link for  
presentation



Link to survey



# Thank you!